

gMark: Schema-Driven Generation of Graphs and Queries

Guillaume Bagan^{*} Angela Bonifati^{*} Radu Ciucanu[†]

George Fletcher[§] Aurélien Lemay[¶] Nicky Advokaat[§]



^{*}Univ Lyon 1 & CNRS LIRIS [†]Univ Clermont Auvergne & CNRS LIMOS [§]TU Eindhoven [¶]Univ Lille 3 & Inria

1. gMark: motivation and goals

Graph-structured data collections are increasingly common in many application domains.

Domain- and application-independent **instance and query workload generators** are important for the experimental study of data-intensive systems.

Currently there is no such framework for graph databases.

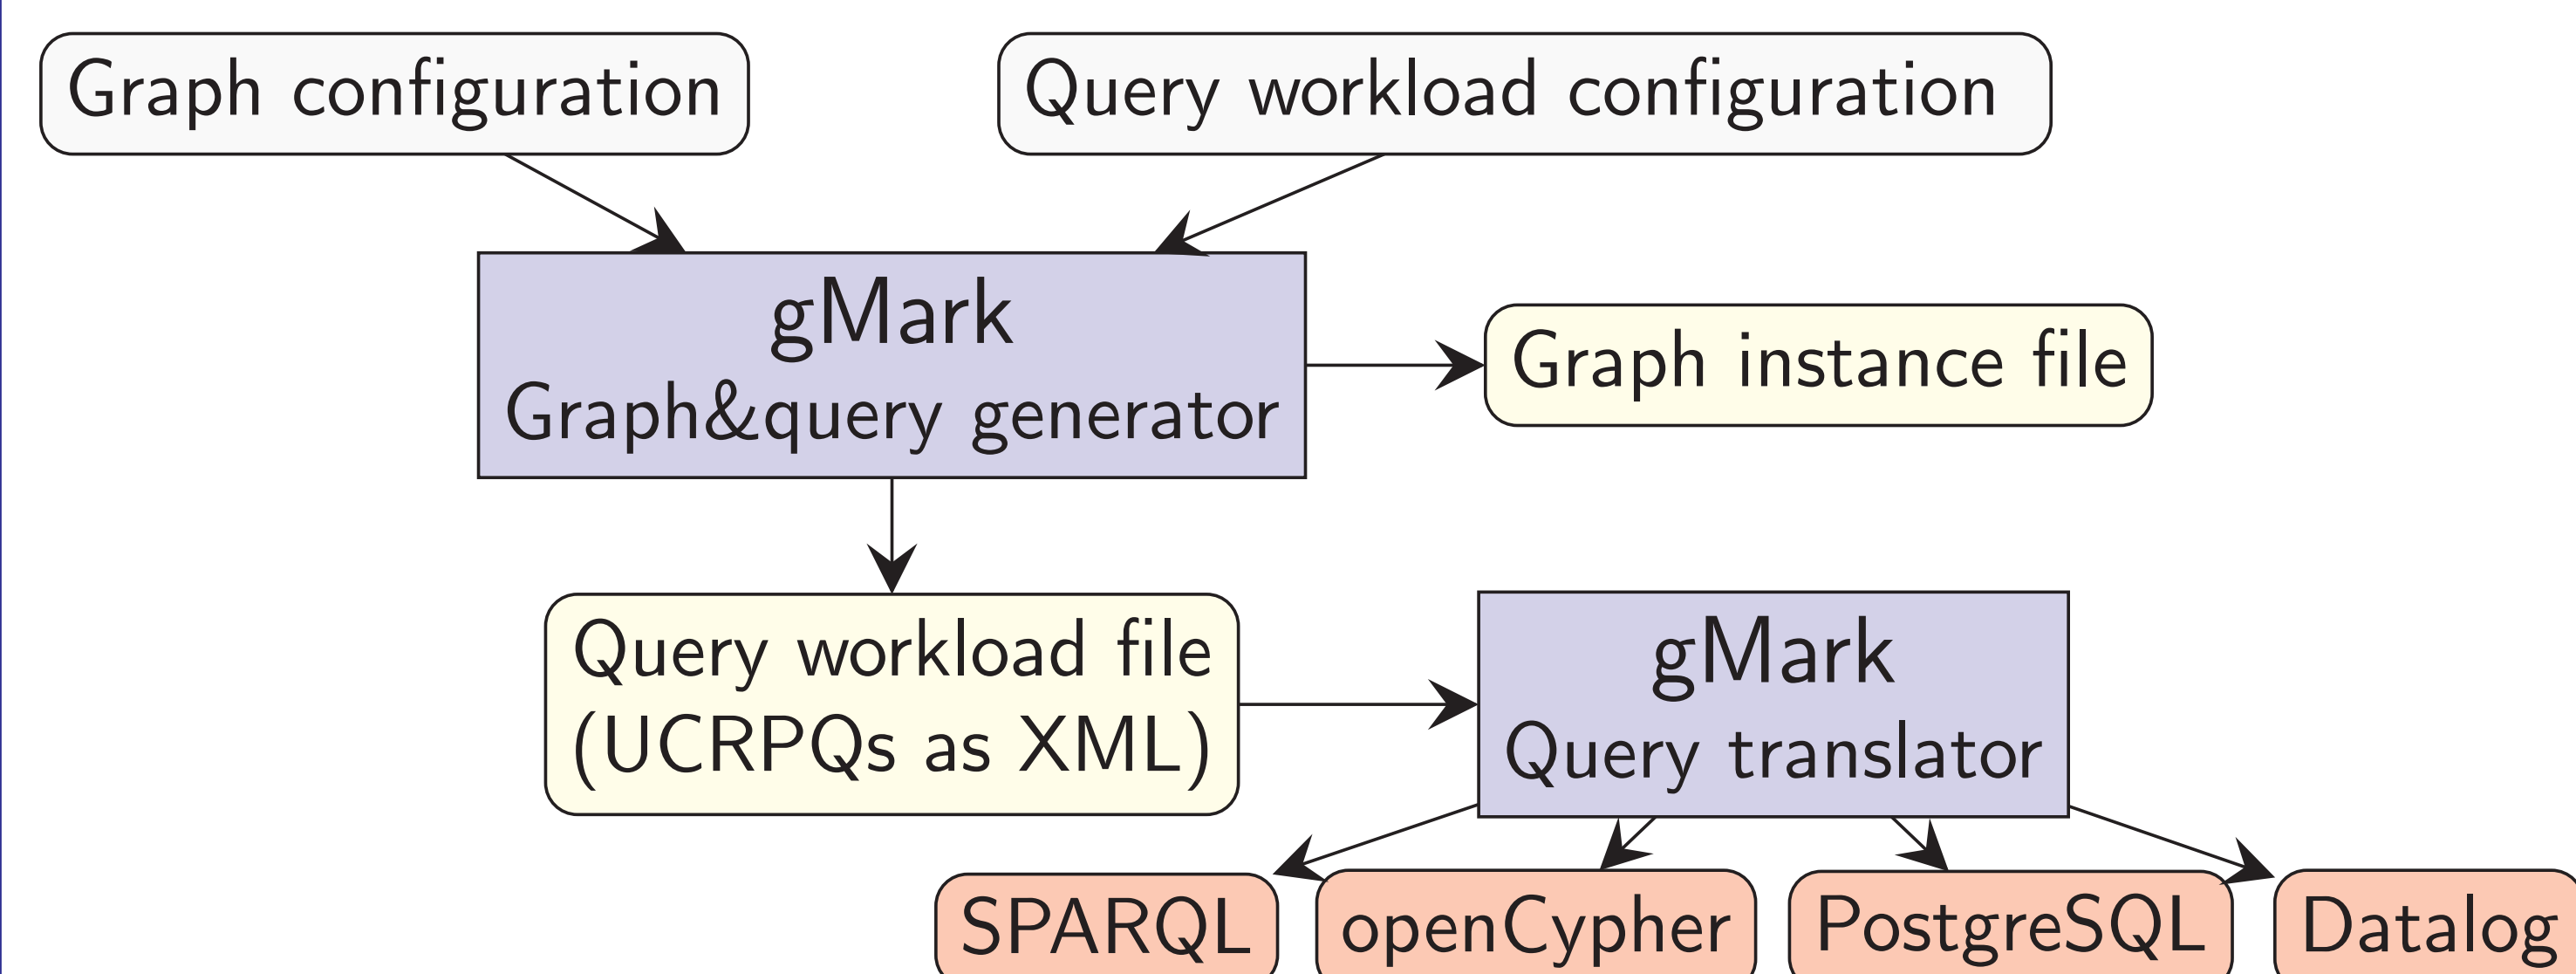
We have developed **gMark**, the first generation framework that

- is schema-driven, providing user-tailored instance and workload diversity, including flexible generation of **recursive path queries**; and
- provides control of query selectivities, in a purely instance-independent schema-driven fashion.

gMark is open-source and ready to use!

<https://github.com/graphMark/gmark>

2. gMark workflow



3. gMark graph configurations

Parameter	Description															
<i>Size</i>	# of nodes to generate															
<i>Node types</i>	finite set of node types ▷ e.g., researcher , paper , conference , city															
<i>Edge predicates</i>	finite set of edge predicates ▷ e.g., authors , publishedIn , heldIn															
<i>Schema constraints</i>	proportion of node types and edge predicates ▷ e.g., 50% of all nodes are researchers															
<i>Degree distributions</i>	on the in- and out-degree of edge predicates ▷ e.g., the number of authors on papers follows a Gaussian distribution, whereas the number of papers authored by a researcher follows a Zipfian															
	<table><tr><td><i>source type</i></td><td><i>predicate</i></td><td><i>target type</i></td><td><i>In-distr.</i></td><td><i>Out-distr.</i></td></tr><tr><td colspan="3">$\xrightarrow{\hspace{1.5cm}}$</td><td></td><td></td></tr><tr><td colspan="3">researcher <u>authors</u> paper</td><td>Gaussian</td><td>Zipfian</td></tr></table>	<i>source type</i>	<i>predicate</i>	<i>target type</i>	<i>In-distr.</i>	<i>Out-distr.</i>	$\xrightarrow{\hspace{1.5cm}}$					researcher <u>authors</u> paper			Gaussian	Zipfian
	<i>source type</i>	<i>predicate</i>	<i>target type</i>	<i>In-distr.</i>	<i>Out-distr.</i>											
$\xrightarrow{\hspace{1.5cm}}$																
researcher <u>authors</u> paper			Gaussian	Zipfian												

4. gMark query workload configurations

Parameter	Description
<i>Workload size</i>	# of queries to generate
<i>Arity</i>	min/max arity of generated queries ▷ e.g., min query arity = 0 and max = 5
<i>Shapes (of join graphs)</i>	chain, star, cycle, and/or star-chain ▷ e.g., generate star and star-chain queries
<i>Selectivity</i>	constant, linear, and/or quadratic ▷ e.g., generate only linear queries
<i>Probability of recursion</i>	percentage of conjuncts having a Kleene star ▷ e.g., 10% occurrence of recursion
<i>Query size</i>	# of conjuncts, # of disjuncts, etc. ▷ e.g., generate only triangles (i.e., queries with exactly 3 conjuncts)

gMark generates **unions of conjunctions of regular path queries** (UCRPQ) e.g., the query $(?x, ?y, ?z) \leftarrow (?x, (a \cdot b + c)^*, ?y), (?y, a, ?w), (?w, b^-, ?z)$ has 3 conjuncts, having 1 or 2 disjuncts, having path length 1 or 2.

5. gMark evaluation

The problems of *graph* and *query generation* are theoretically **intractable** \Rightarrow We developed *best-effort* algorithms that work in *linear time*.

(a) gMark scalability

Instance generation. We adapted the scenarios of several popular use cases into meaningful gMark configurations:

- **Bib**: our default bibliographical use-case
- **LSN**: LDBC social network benchmark
- **WD**: WatDiv e-commerce benchmark
- **SP**: SP2Bench DBLP benchmark

	100K	1M	10M	100M
Bib	0m0.057s	0m0.638s	0m8.344s	1m28.725s
LSN	0m0.225s	0m1.451s	0m23.018s	3m11.318s
WD	0m2.163s	0m25.032s	4m10.988s	113m31.078s
SP	0m0.638s	0m7.048s	1m28.831s	15m23.542s

Graph instance generation times, with varying graph sizes (# nodes)

Generation time depends heavily on density of instances (e.g., WD has 100x number of edges than Bib).

Query generation. gMark *generates* workloads of *1000 queries* for Bib in $\sim 0.3s$; LSN and SP in $\sim 1.5s$; and, for the richer WD scenario in $\sim 10s$.

Query translation of the 1000 queries into all four supported syntaxes for each of the four scenarios required $\sim 0.1s$.

(b) gMark accuracy for the query selectivity estimation

Given a binary query Q and a graph G , we assume that $|Q(G)| = |G|^\alpha$, where α is the **selectivity value** (0–constant, 1–linear, 2–quadratic).

- Experiments confirmed the assumption and the estimation quality.


6. Evaluation of current graph DBMSs with gMark

We studied **query evaluation performace** of four mainstream graph DBMSs:

- **P**: PostgreSQL v9.3.9 (SQL:1999 recursive views)
- **S**: a popular SPARQL query engine (SPARQL 1.1)
- **G**: a native graph database (openCypher)
- **D**: a modern Datalog engine (Datalog)

(a) Non-recursive queries

Query execution times for various graph sizes and diverse query workloads: *Len* (varying path lengths, 1 disjunct, 1 conjunct), *Dis* (multiple disjuncts, 1 conjunct), *Con* (multiple conjuncts and disjuncts):




Constant queries

Linear queries

Quadratic queries

(b) Recursive queries

Query execution times for simple recursive queries on various small graph sizes (from 2K to 32K nodes):



References

- [1] G. Bagan, A. Bonifati, R. Ciucanu, G. Fletcher, A. Lemay, N. Advokaat. **gMark: Schema-Driven Generation of Graphs and Queries**. IEEE TKDE, 2017. <http://arxiv.org/abs/1511.08386>
- [2] G. Bagan, A. Bonifati, R. Ciucanu, G. Fletcher, A. Lemay, N. Advokaat. **Generating flexible workloads for graph databases**. VLDB Demo, 2016.